

# CIFAR-10 Classification Problem

Zhilin Fan

Professor: John Cunningham  
GR5242 Advanced Machine Learning

December 22, 2017

## Abstract

In this project, we built up a Convolutional Neural Network(CNN) to solve the classification problem on CIFAR-10 datasets, which consists of 60000  $32 \times 32$  colour images in 10 classes, with 6000 images per class. And there are 50000 training images and 10000 test images. We began with a simple CNN model but met with a serious overfitting problem. The test accuracy was between 50% – 55% while the training loss was about 0.2 and training accuracy was about 90%. We tried several ways, including adding more layers, adding dropout layer, changing the activation function of the layer, but these did not make big difference. Finally, we used the following strategies:

1. preprocess the image data(crop image from  $32 \times 32 \times 3$  to  $28 \times 28 \times 3$  and add some noise in the training step
2. add weight decay in two fully-connected layers

Finally, we improved the performance and the accuracy was about 67%.

In addition, we recomputed the bottleneck and retrained the last layer of Inception-V3 model. The final test accuracy was about 81.5%.

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>CNN model</b>	<b>3</b>
2.1	Preprocess the data . . . . .	3
2.2	Relu activation . . . . .	3
2.3	Weight Decay and L2 Regularization . . . . .	5
2.4	Summary and Perspective . . . . .	5
<b>3</b>	<b>Inception-V3 Model</b>	<b>6</b>
<b>4</b>	<b>Conclusion</b>	<b>7</b>
	<b>Appendices</b>	<b>8</b>
<b>A</b>	<b>Python Code</b>	<b>8</b>

# 1 Introduction

CIFAR-10 is a really important image classification dataset for data scientists and machine learning researchers. In the past few decades a large number of projects are finished based on CIFAR-10. In this project we are going to create our deep learning classifier for CIFAR-10 with the highest accuracy.

## 2 CNN model

In this section, we implemented a 9-layer Convolutional Neural Network to solve the CIFAR-10 classification problem. The construction of this network is as following:

1. a convolutional layer: filter size [3, 5, 5, 32], Relu activated
2. a max-pooling layer: kernel size [2, 2] and stride 2
3. a convolutional layer: filter size [32, 5, 5, 64], Relu activated
4. a max-pooling layer: kernel size [2, 2] and stride 2
5. a convolutional layer: filter size [64, 5, 5, 64], Relu activated
6. a convolutional layer: filter size [64, 3, 3, 32], Relu activated
7. a max-pooling layer: kernel size [3, 3] and stride 2
8. a fully-connected layer: filter size [3 \* 3 \* 32, 384], Relu activated, standard deviation=0.04, weight decay=0.001, biases=0.1
9. a fully-connected layer: filter size [384, 192], Relu activated, standard deviation=0.04, weight decay=0.01, biases=0.1

The following graph shows the structure of this model:

At the beginning, we trained the network without any preprocessing or weight decay, and met with serious overfitting problem. Several strategies are used to avoid this problem.

### 2.1 Preprocess the data

Adding some noise to the initial data can improve the stationary and test performance of the model. Thus, we firstly cropped all the image data from [32, 32, 3] to [28, 28, 3] since the margin of a image often contains little information. Then, for the training data, we randomly cropped the image, flipped the original image horizontally, changed the brightness and the contrast of the image. For test images, we only cropped the center of the image.

After applying this process, the loss decreased much more slowly than before and the final loss was larger. The test performance was getting better, indicating the overfitting problem was partly solved. However, we still cannot get the accuracy larger than 60%.

### 2.2 Relu activation

We used Relu activation in every layer because it has advantages in unsupervised neural networks. In *Deep Sparse Rectifier Neural Networks*(Xavier, 2011), it has been proved that while logistic sigmoid neurons are more biologically plausible than hyperbolic tangent neurons, the latter work better for training multi-layer neural networks. However, rectifying neurons are an even better model of biological neurons and yield equal or better performance than hyperbolic tangent networks in spite of the hard non-linearity and non-differentiability at zero, creating sparse representations with true zeros, which seem remarkably suitable for naturally sparse data. Even though they can take advantage of semi-supervised setups with extra-unlabeled data, deep rectifier networks can reach their best performance without requiring any unsupervised pre-training on purely supervised tasks with large labeled datasets.[1]

## Main Graph

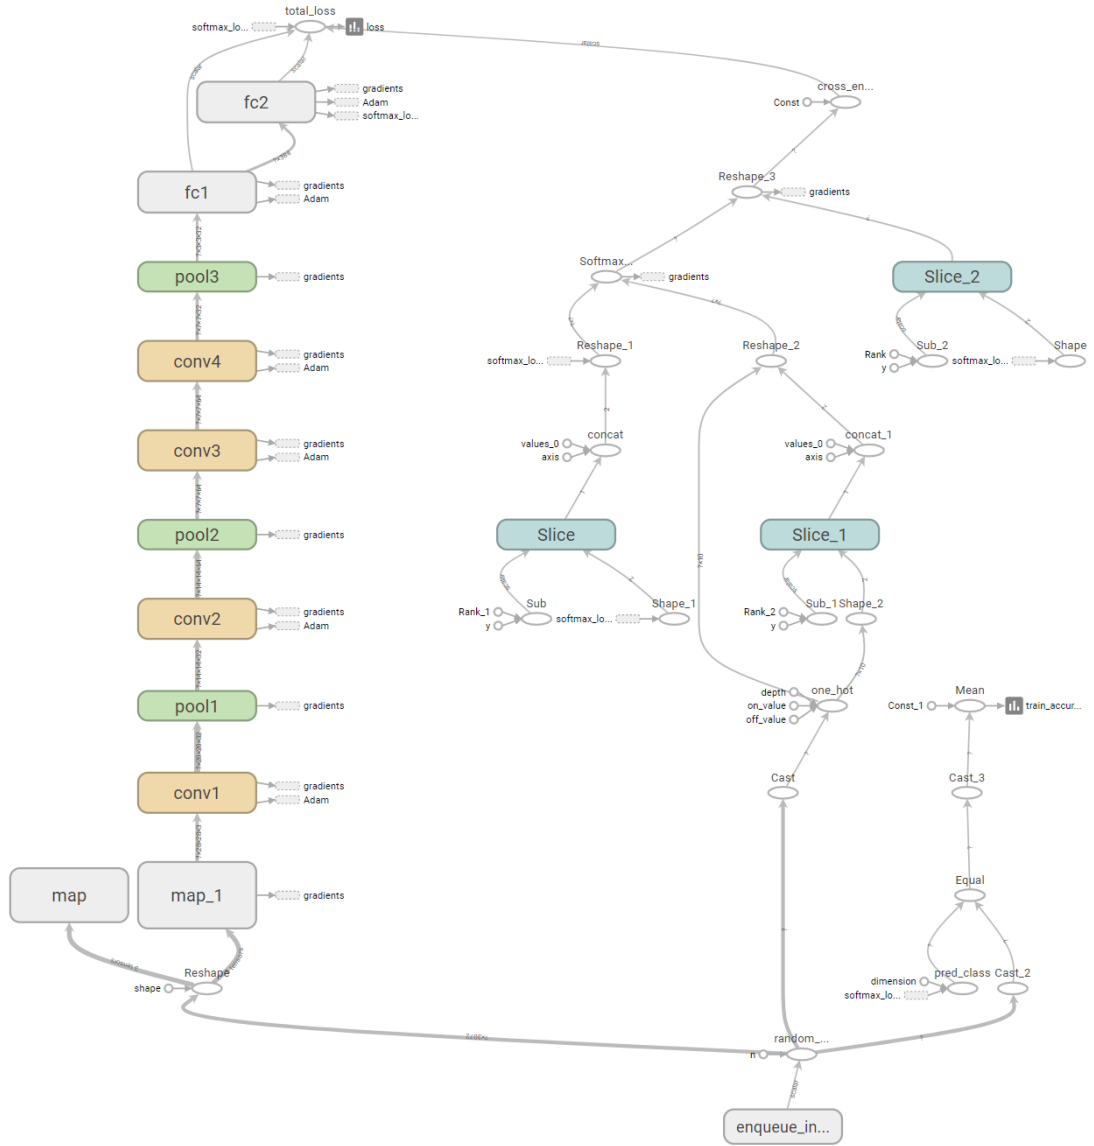


Figure 1: CNN model

## 2.3 Weight Decay and L2 Regularization

To avoid overfitting, we used L2 Regularization in our network.

In machine learning, overfitting problems often caused by too many features, so we need to reduce features or add penalty weights to unimportant features. And regularization can help add penalty, that is, the total loss of the model would contain the feature weights. Then it can filter the most efficient features and avoid overfitting.

Generally, L1 Regularization would set most feature weights to 0 and produces sparse coefficients while L2 Regularization would produces more average weights.

Table 1: Difference between L2 and L1

L2		L1	
L2 Loss Function	L2 Regularization	L2 Loss Function	L2 Regularization
Not very robust	Computational efficient	Robust	Computational inefficient
Stable solution	Non-sparse outputs	Unstable solution	Sparse outputs
Always one solution	No feature selection	Possibility multiple solution	Built-in feature selection

## 2.4 Summary and Perspective

Finally, we implemented 38000+ steps after these adjustments, the final performance was satisfying and the following graphs 2, 3, 4 show the training loss and accuracy and a screen shot of the test accuracy at 34000<sub>th</sub> step.




Figure 2: Training Loss




Figure 3: Training Accuracy

```

INFO:tensorflow:loss = 0.720076, step = 33601 (82.169 sec)
INFO:tensorflow:global_step/sec: 1.21349
INFO:tensorflow: (82.368 sec)
INFO:tensorflow:loss = 0.782067, step = 33601 (82.367 sec)
INFO:tensorflow:global_step/sec: 1.20695
INFO:tensorflow: (82.874 sec)
INFO:tensorflow:loss = 0.718546, step = 33701 (82.872 sec)
INFO:tensorflow:global_step/sec: 1.20888
INFO:tensorflow: (82.702 sec)
INFO:tensorflow:loss = 0.730189, step = 33801 (82.705 sec)
INFO:tensorflow:global_step/sec: 1.19948
INFO:tensorflow: (83.399 sec)
INFO:tensorflow:loss = 0.717277, step = 33901 (83.398 sec)
INFO:tensorflow:Saving checkpoints for 34000 into ./outhook/convnet_model\model.ckpt.
INFO:tensorflow:Loss for final step: 0.788978.
INFO:tensorflow:Starting evaluation at 2017-12-12-19:41:44
INFO:tensorflow:Restoring parameters from ./outhook/convnet_model\model.ckpt-34000
INFO:tensorflow:Finished evaluation at 2017-12-12-19:42:09
INFO:tensorflow:Saving dict for global step 34000: global_step = 34000, loss = 0.983667, test_accuracy = 0.674
('loss': 0.98366714, 'test_accuracy': 0.67400002, 'global_step': 34000)

```

Figure 4: Screen Shot

From the plot, we can see that the accuracy increased very slowly after 20k steps. The final training accuracy is about 75%, close to the final test accuracy 67.4%, indicates the model is not overfitting .

To further improve the test performance, we could add more layers or use some existed feature extraction process.

### 3 Inception-V3 Model

In this part, we use the similar strategy in hw4, but we cropped the image to  $28 \times 28$  then reshaped it to  $299 \times 299$ , recomputed the bottlenecks and got the final training accuracy after 2500 steps was about 90.2%, the test accuracy was about 81.5%.

The following graphs show the accuracy and cross entropy of the training step and a screen shot of the detailed value of train and test accuracy.




Figure 5: Inception-V3 ac and loss

---

```

Accuracy at step 800 is 0.8669999837875366
Accuracy at step 900 is 0.8703600168228149
Accuracy at step 1000 is 0.872979998588562
Accuracy at step 1100 is 0.8761399984359741
Accuracy at step 1200 is 0.8786799907684326
Accuracy at step 1300 is 0.8813999891281128
Accuracy at step 1400 is 0.8837199807167053
Accuracy at step 1500 is 0.8861200213432312
Accuracy at step 1600 is 0.8881000280380249
Accuracy at step 1700 is 0.8899999856948853
Accuracy at step 1800 is 0.8920400142669678
Accuracy at step 1900 is 0.8938800096511841
Accuracy at step 2000 is 0.8953400254249573
Accuracy at step 2100 is 0.897159993648529
Accuracy at step 2200 is 0.8986799716949463
Accuracy at step 2300 is 0.9002799987792969
Accuracy at step 2400 is 0.9016799926757812
----- Training done! -----
966 9967 9968 9969 9970 9971 9972 9973 9974 9975 9976 9977 9978 9979 9980 9981 9982 9983 9984 9985 9986 9987 9988 9989 9990 9991 9992 9993
9994 9995 9996 9997 9998 9999 Evaluation accuracy was: 0.8148999810218811

```

---

Figure 6: Inception-V3 training and test

## 4 Conclusion

We already showed two models we worked on for CIFAR-10 classification problem: CNN model and Inception-V3 model. As we can see from the result, Inception-V3 model, without overfitting problem, gives a higher accuracy on both training set and testing set than the 9-layer CNN model.

Indeed, the advanced magical way Inception-V3 model used to extract the features is the key to the good performance. So it is really important to pay attention to feature extraction process during the design of our own CNN.

## References

- [1] Glorot, Xavier and Bordes, Antoine and Bengio, Yoshua. *Deep Sparse Rectifier Neural Networks*, Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics, 2011, pp315-323.

# Appendices

## A Python Code

Part 1: CNN model

*# import the package we need*

**import** numpy as np

**import** tensorflow as tf

**def** unpickle(**file**):

*'''Read the raw data, a dictionary which contains the features,  
labels and other information'''*

**import** pickle

**with** **open**(**file**, 'rb') as fo:

**dict** = pickle.load(fo, encoding='bytes')

**return dict**

**def** get\_train():

*'''Get all the 5 trainsets and convert the datatype Unit8 to float32*

*Output:*

*trdata: [50000, 3072] RGB features*

*trlables: [50000,] training labels from 0-9*

*'''*

**data** = unpickle('./cifar-10-batches-py/data\_batch\_1')

**trdata** = np.asarray(**data**[b'data'], dtype= np.float32)

**trlables** = np.asarray(**data**[b'labels'], dtype= np.float32)

**for** **i** **in** **range**(2,6):

**data** = unpickle('./cifar-10-batches-py/data\_batch\_{}'.format(**i**))

**trdata** = np.row\_stack((**trdata**, (np.asarray(**data**[b'data'], dtype= np.float32))))

**trlables** = np.concatenate((**trlables**, (np.asarray(**data**[b'labels'],  
dtype= np.float32))), 0)

**return**(**trdata**, **trlables**)

**def** get\_test():

*'''Get the testsets and convert the datatype Unit8 to float32*

*Output:*

*tedata: [10000, 3072] RGB features*

*telables: [10000,] test labels from 0-9*

*'''*

**data**= unpickle('./cifar-10-batches-py/test\_batch')

**tedata** = np.asarray(**data**[b'data'], dtype= np.float32)

**telables** = np.asarray(**data**[b'labels'], dtype= np.float32)

**return**(**tedata**, **telables**)

*# add the weight decay to avoid overfitting*

**def** weight\_decay\_variable(name, shape, stddev, wd):

**dtype** = tf.float32

**var** = tf.get\_variable(name, shape,

**initializer**=tf.truncated\_normal\_initializer(stddev=stddev),

**dtype**=dtype)

*# define the weight decay*

**weight\_decay** = tf.multiply(tf.nn.l2\_loss(**var**), wd, name='weight\_loss')



```

    # add the weight to total loss
    tf.add_to_collection('losses', weight_decay)
    return var

crop_size = 28

# for training step, crop and distort the image
def preprocess_train_image(x):
    '''Input:
        x: a training data with size [1, 3072]
    Output:
        float_image: a distorted training data with size [1, 28x28x3]
    '''
    # Randomly crop the original image.
    distorted_image = tf.random_crop(x, [crop_size, crop_size, 3])

    # Randomly flip the original image horizontally.
    distorted_image = tf.image.random_flip_left_right(distorted_image)

    # Randomly distort the brightness of the image.
    distorted_image = tf.image.random_brightness(distorted_image, max_delta=1.0)

    # Randomly distort the contrast of the image.
    distorted_image = tf.image.random_contrast(distorted_image, lower=0.2, upper=1.0)

    # Subtract off the mean and divide by the variance of the pixels.
    float_image = tf.image.per_image_standardization(distorted_image)

    # Set the dimension of the output
    float_image.set_shape([crop_size, crop_size, 3])
    return (float_image)

# for testing step, only crop the image
def preprocess_test_image(x):
    '''Input:
        x: a testing data with size [1, 32*32*3]
    Output:
        float_image: a cropped data with size [1, 28*28*3]
    '''
    # crop the center of the image
    resized_image = tf.image.resize_image_with_crop_or_pad(x, crop_size, crop_size)
    float_image = tf.image.per_image_standardization(resized_image)
    float_image.set_shape([crop_size, crop_size, 3])
    return(float_image)

def create_cnnmodel(features, labels, mode):
    """Build the CNN Model"""
    # Input layer
    input_features = tf.reshape(features["x"], [-1, 32, 32, 3])
    input_layer = tf.map_fn(preprocess_test_image, input_features)
    if mode == tf.estimator.ModeKeys.TRAIN:
        input_layer = tf.map_fn(preprocess_train_image, input_features)

```

```

# Convolutional Layer #1 (output:[batch_size , 28, 28, 32])
conv1 = tf.layers.conv2d(inputs=input_layer , filters=32, kernel_size=[5, 5],
                        padding="same", activation=tf.nn.relu , name="conv1")

# Pooling Layer #1 (output:[batch_size , 14, 14, 32])
pool1 = tf.layers.max_pooling2d(inputs=conv1 , pool_size=[2, 2],
                                strides=2, name="pool1")

# Convolutional Layer #2 (output:[batch_size , 14, 14, 64])
conv2 = tf.layers.conv2d(inputs=pool1 , filters= 64, kernel_size=[5, 5],
                        padding="same", activation=tf.nn.relu , name="conv2")

# Pooling Layer #2 (output:[batch_size , 7, 7, 64])
pool2 = tf.layers.max_pooling2d(inputs=conv2 , pool_size=[2, 2],
                                strides=2, name="pool2")

# Convolutional Layer #3 (output:[batch_size , 7, 7, 64])
conv3= tf.layers.conv2d(inputs=pool2 , filters= 64, kernel_size=[5, 5],
                        padding="same", activation=tf.nn.relu , name="conv3")

# Convolutional Layer #4 (output:[batch_size , 7, 7, 32])
conv4 = tf.layers.conv2d(inputs=conv3 , filters= 32, kernel_size=[3, 3],
                        padding="same", activation=tf.nn.relu , name="conv4")

# Pooling Layer #3 (output:[batch_size , 3, 3, 32])
pool3 = tf.layers.max_pooling2d(inputs=conv4 , pool_size=[3, 3],
                                strides=2, name="pool3")

# Fully-connected Layer #1 (output:[batch_size , 384])
with tf.variable_scope('fc1') as scope:
    reshape = tf.reshape(pool3 , shape=[-1, 3*3*32])
    dim = reshape.get_shape()[1].value
    weights = weight_decay_variable("weights_fc1" , shape=[dim, 384],
                                    stddev=0.04, wd=0.001)
    biases = tf.get_variable("biases_fc1" , shape=[384],
                             initializer=tf.constant_initializer(0.1))
    fc1 = tf.nn.relu(tf.matmul(reshape , weights)+ biases , name=scope.name)

# Fully-connected Layer #2 (output:[batch_size , 192])
with tf.variable_scope("fc2") as scope:
    weights = weight_decay_variable("weights_fc2" , shape=[384, 192],
                                    stddev=0.04, wd=0.01)
    biases = tf.get_variable("bias_fc2" , shape=[192],
                             initializer=tf.constant_initializer(0.1))
    fc2 = tf.nn.relu(tf.matmul(fc1 , weights)+ biases , name=scope.name)

# Logits Layer (output:[batch_size , 10])
with tf.variable_scope("softmax_logits") as scope:
    weights = weight_decay_variable("weights_out" , shape=[192, 10],
                                    stddev=1.0/192, wd=0.0)
    biases = tf.get_variable("biases_out" , shape=[10])
    logits = tf.add(tf.matmul(fc2 , weights), biases , name=scope.name)

```

```

# Prediction dictionary: pred_class is the prediction labels , from 0 to 9.
predictions = {
    "pred_class": tf.argmax(logits , axis=1, name= "pred_class"),
    "probabilities": tf.nn.softmax(logits , name="softmax_tensor")
}

if mode == tf.estimator.ModeKeys.PREDICT:
    return tf.estimator.EstimatorSpec(mode=mode, predictions=predictions)

# Calculate cross entropy and the total losses of each batch (for training steps)
# change the labels into one_hot type, depth=10 means 10 labels
onehot_labels = tf.one_hot(indices=tf.cast(labels , tf.int32), depth=10)
# compute the cross entropy between the logits and true labels
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels = onehot_labels ,
                                                         logits= logits)

# compute the cross entropy loss
cross_entropy_loss = tf.reduce_mean(cross_entropy , name='cross_entropy_loss')
tf.add_to_collection('losses' , cross_entropy_loss)
# total loss is the cross entropy loss and the weight decay loss
loss = tf.add_n(tf.get_collection('losses') , name='total_loss')
# compute the training accuracy for visualization in the tensorboard
labels_int = tf.cast(labels , tf.int64)
accuracy = tf.reduce_mean(tf.cast(tf.equal(predictions["pred_class"] ,
                                           labels_int) , tf.float32))

# Construct the training optimizer
if mode == tf.estimator.ModeKeys.TRAIN:
    # using AdamOptimizer and the learning rate = 0.001
    global_step= tf.train.get_global_step()
    optimizer = tf.train.AdamOptimizer(learning_rate=0.001)
    train_op = optimizer.minimize( loss=loss , global_step=tf.train.get_global_step())

    # Create the summary graph for loss , logits and training accuracy
    #cd C:\Users\qn_li\Columbia\Advanced Machine Learning\Project
    #tensorboard --logdir=./outhook/tb
    tf.summary.scalar("loss" , loss)
    tf.summary.histogram('logits' , logits)
    tf.summary.scalar('train-accuracy' , accuracy)
    summary_hook = tf.train.SummarySaverHook(save_steps=1, output_dir='./outhook/tb' ,
                                              summary_op=tf.summary.merge_all())

    return tf.estimator.EstimatorSpec(mode=mode, loss=loss , train_op=train_op , training_ops=summary_hook)

# Compute the test accuracy and construct the evaluate optimizer
eval_metric_ops = { "test_accuracy": tf.metrics.accuracy(labels=labels , predictions=predictions) }

return tf.estimator.EstimatorSpec(
    mode=mode, loss=loss , eval_metric_ops=eval_metric_ops)

def main():
    # Load training and eval data and convert the datatype from Unit8 to float32
    train_data , train_labels = get_train()
    eval_data , eval_labels = get_test()

```

```

# Create the Estimator
cifar10_classifier = tf.estimator.Estimator(model_fn= create_cnnmodel , model_dir="./ou

# Set up logging for predictions
#tensors_to_log = {"probabilities": "softmax_tensor"}
tensors_to_log={}
logging_hook = tf.train.LoggingTensorHook( tensors=tensors_to_log , every_n_iter=100)

# Train the model
train_input_fn = tf.estimator.inputs.numpy_input_fn(
    x={"x": train_data}, y=train_labels ,
batch_size= 128, num_epochs=100, shuffle=True)

train_results = cifar10_classifier.train( input_fn=train_input_fn , steps= total_steps ,

# Evaluate the model and print results
eval_input_fn = tf.estimator.inputs.numpy_input_fn(x={"x": eval_data},
                                                    y=eval_labels ,
                                                    num_epochs=1, shuffle=False)

eval_results = cifar10_classifier.evaluate(input_fn=eval_input_fn)
print(eval_results)

```

```
total_steps = 38000
```

```
main()
```

Part 2: Inception\_V3 model

```

import tensorflow as tf
import numpy as np
import os
os.environ[ 'TF_CPP_MIN_LOG_LEVEL' ]= '2'
import os.path
import transfer_learning_project as transfer_learning
import matplotlib.pyplot as plt
import scipy.misc
import glob
import sys
import random

```

```

# Ensure target log dir exists
INCEPTION_LOG_DIR = './tmp/inception_v3_log'
if not os.path.exists(INCEPTION_LOG_DIR):
    os.makedirs(INCEPTION_LOG_DIR)

```

```

# import training and test data
training_images , _, label_maps = transfer_learning.create_image_lists(
    './cifar-10/train' , testing_percentage=0, max_number_images=5001)

```

```

testing_images , _, _ = transfer_learning.create_image_lists(
    './cifar-10/test' , testing_percentage=0, max_number_images=1001)

```

```
# check the length of the training_images and testing_images
```

```

print(len(training_images))
print(len(testing_images))

# Create the inception model.
graph, bottleneck, resized_input, softmax = transfer_learning.create_model()

# Use a summary writer to write the loaded model graph and display it in tensorboard.
# tensorboard --logdir=./tmp/inception_v3-log
with graph.as_default():
    jpeg_data, decoded_image = transfer_learning.make_jpeg_decoding()
    # Define summaries for tensorboard
    # create histogram summary for output softmax and bottleneck
    tf.summary.histogram('output', softmax)
    tf.summary.histogram('bottleneck', bottleneck)
    # create summary for input image
    tf.summary.image('input', resized_input)
    summary_op = tf.summary.merge_all()

    with tf.Session() as sess:
        summary_writer = tf.summary.FileWriter(INCEPTION_LOG_DIR, sess.graph)
        sess.run(tf.global_variables_initializer())

def classify_image(session, image_path, summary_op):
    """This functions reads a single image from disk and classifies
    the image using the pre-trained network.

    Parameters
    

---


    session: the tensorflow session to use to run the operation
    image_path: a string corresponding to the path of the image to load
    summary_op: the summary operation.

    Returns
    

---


    label: an integer representing the label of the classified example.
    softmax_output: the network's output multinomial probabilities
    """

    # Open single image file and extract data
    with open(image_path, 'rb') as f:
        image_data = f.read()
    # run the train step and add the summary to tensorboard
    input_value = session.run((decoded_image), {jpeg_data: image_data})
    softmax_output = session.run(softmax, feed_dict={resized_input: input_value})
    summary = session.run(summary_op, {jpeg_data: image_data, resized_input: input_value})
    summary_writer = tf.summary.FileWriter(INCEPTION_LOG_DIR, session.graph)
    summary_writer.add_summary(summary)
    # Return label
    return(np.argmax(softmax_output), softmax_output)

def compute_bottleneck(session, image_data):
    """Computes the bottleneck for a given image

    Parameters

```

---

```

session: the tensorflow session to use for the computation.
jpeg_data_tensor: the tensor to feed the jpeg data into.
bottleneck_tensor: the tensor representing.
image_data: a byte sequence representing the encoded jpeg data.

Returns

```

---

```

A numpy array containing the bottleneck information for the image.
"""

new_value = session.run(decoded_image, feed_dict= {jpeg_data: image_data})
bottleneck_values = session.run(bottleneck, feed_dict= {resized_input: new_value})
return(bottleneck_values)

# This cell generates all the bottlenecks and it will take a
# long time since there are 50000 bottlenecks should be computed.
with graph.as_default():
    with tf.Session() as session:
        transfer_learning.cache_bottlenecks(compute_bottleneck, session, training_images)

# This loads the training data as a matrix of training examples
# and a vector of labels
training_data_set = transfer_learning.create_training_dataset(training_images)

def make_final_layers(bottleneck_tensor, num_classes):
    """Create the operations for the last layer of the network to be retrained.
    Parameters

```

---

```

bottleneck_tensor: the bottleneck tensor in the original network
num_classes: the number of output classes

Returns

```

---

```

bottleneck_input: the input placeholder for the bottleneck values
label_input: the input placeholder for the label values
logits: the tensor representing the unnormalized log probabilities
for each class.
train_step: an operation representing one gradient descent step.
"""

bottleneck_tensor_size = int(bottleneck.shape[1])

with tf.variable_scope('input'):
    # This is the input for the bottleneck. It is created
    # as a placeholder with default. During training, we will
    # be passing in the bottlenecks, but during evaluation,
    # the value will be propagated from the bottleneck computed
    # from the image using the full network.
    bottleneck_input = tf.placeholder-with-default(
        bottleneck_tensor,
        [None, bottleneck_tensor_size],
        'bottleneck_input')

    # This is the input for the label (integer, 1 to number of classes)

```

```

        label_input = tf.placeholder(tf.int64, [None], name='label_input')

# Define weights, biases, and logit transforms
logits = tf.layers.dense(bottleneck_input, num_classes)
# Compute the cross entropy loss
loss = tf.losses.sparse_softmax_cross_entropy(labels=label_input, logits=logits)
# Create a summary for the loss
loss_summary = tf.summary.scalar('cross_entropy', loss)
# Create a Gradient Descent Optimizer
# optimizer = tf.train.GradientDescentOptimizer(0.1)
optimizer = tf.train.AdamOptimizer(learning_rate=0.001)
# Obtain a function which performs a single training step
train_step = optimizer.minimize(loss)
return bottleneck_input, label_input, logits, train_step, loss_summary

def compute_accuracy(labels, logits):
    """Compute the accuracy for the predicted output.

    Parameters
    -----
    labels: The input labels (in a one-hot encoded fashion).
    predicted_output: The predicted class probability for each output.

    Returns
    -----
    A tensor representing the accuracy.
    """
    prediction = tf.argmax(logits, 1, name='pred_class')
    accuracy = tf.reduce_mean(tf.cast(tf.equal(prediction, labels), tf.float32))
    accuracy_summary = tf.summary.scalar('accuracy', accuracy)

    return accuracy, accuracy_summary

# We create the necessary operations to fine tune the model.

with graph.as_default():
    bottleneck_input, label_input, logits, train_step, loss_summary = make_final_layers(bottleneck_input, label_input, num_classes)
    accuracy, accuracy_summary = compute_accuracy(label_input, logits)
    summary_op = tf.summary.merge([loss_summary, accuracy_summary])

def execute_train_step(session: tf.Session, summary_writer: tf.summary.FileWriter, current_step: int):
    """This function runs a single training step.

    You may wish to print some progress information as you go along.

    Parameters
    -----
    session: the tensorflow session to use to run the training step.
    summary_writer: the summary file writer to write your summaries to.
    current_step: the current step count (starting from zero)
    """
    _, ac, summary = session.run((train_step, accuracy, summary_op),
                                feed_dict={bottleneck_input: training_data_set['bottlenecks'],

```

```

        label_input: training_data_set['labels']
    })

summary_writer.add_summary(summary, current_step)

if current_step % 100 == 0:
    print('Accuracy at step {0} is {1}'.format(current_step, ac))

def evaluate_images(session: tf.Session, images_jpeg_data: [bytes], labels: [int]):
    """This function will evaluate the accuracy of our model on the specified data.

    Parameters
    -----
    session: the tensorflow session to use to run the evaluation step.
    images_jpeg_data: a list of strings, with each element in the list corresponding
        to the jpeg-encoded data for a given image
    labels: a list of integers, with each element in the list corresponding to the label
        of a given image.

    Returns
    -----
    This function should return the accuracy as a floating point number between
    0 and 1 (proportion of correctly classified instances).
    """
    correct = []
    i = 0
    for label, jpeg in zip(labels, images_jpeg_data):
        image_data = session.run(decoded_image, feed_dict={jpeg_data: jpeg})
        ac = session.run(accuracy, feed_dict={resized_input: image_data, label_input: [label]})
        correct.append(ac)
        print(i, end=" ")
        i+=1

    return np.mean(correct)

# run the training and evaluation!
with graph.as_default():
    with tf.Session() as session:
        print('_____Starting training_____')
        session.run(tf.global_variables_initializer())
        summary_writer = tf.summary.FileWriter(os.path.join(INCEPTION_LOG_DIR, 'retrained'))
        for i in range(2500):
            execute_train_step(session, summary_writer, i)
        summary_writer.close()
        print('_____Training done!_____')
        print('_____Loading testing data_____')
        tlabels, timages = transfer_learning.get_testing_data(testing_images)
        print('_____Evaluating on testing_____')

        eval_accuracy = evaluate_images(session, timages, tlabels)
        print('Evaluation accuracy was: {0}'.format(eval_accuracy))

```